

```
In [ ]: #!/pip install vaderSentiment
#!/pip uninstall scikit-learn -y
#!/pip install -U scikit-learn
#!/pip install gensim

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from google.colab import drive
from collections import Counter

import re
import nltk
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from wordcloud import WordCloud
import matplotlib.pyplot as plt
from gensim import corpora
from gensim.models import LdaModel
from gensim.models.coherencemodel import CoherenceModel
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
nltk.download('stopwords')
nltk.download('punkt')

from sklearn.cluster import KMeans
from scipy.spatial.distance import pdist, squareform
from scipy.cluster.hierarchy import linkage, dendrogram
from scipy.cluster.hierarchy import fcluster
from sklearn.cluster import DBSCAN
from sklearn.neighbors import NearestNeighbors
from sklearn.metrics import silhouette_score

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree, export_text
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.metrics import RocCurveDisplay
from sklearn.ensemble import BaggingClassifier, AdaBoostClassifier, RandomForestClassifier
from sklearn.model_selection import ValidationCurveDisplay, validation_curve
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...  
[nltk_data] Package stopwords is already up-to-date!  
[nltk_data] Downloading package punkt to /root/nltk_data...  
[nltk_data] Package punkt is already up-to-date!
```

```
In [ ]: drive.mount('/content/drive')  
df = pd.read_csv('/content/drive/My Drive/Colab Notebooks/INF03237_Spring2024/HW/Health_Data.csv', na_values = ["N/A"])  
  
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

```
In [ ]: # Dataset size  
df.shape
```

```
Out[ ]: (3460, 10)
```

```
In [ ]: # Drop missing values  
df.dropna(inplace=True)  
df.shape
```

```
Out[ ]: (3452, 10)
```

```
In [ ]: # Summary statistics  
df.describe()  
df.dtypes
```

```
Out[ ]: Reddit_ID      object  
Subreddits      object  
Titles          object  
Body            object  
Author          object  
Initial Create  object  
Retrieved_date  object  
Date_of_collection  object  
Number of Comments  float64  
Decision        float64  
dtype: object
```

```
In [ ]: df.head()
```

Out []:

	Reddit_ID	Subreddits	Titles	Body	Author	Initial Create	Retrieved_date	Date_of_collection	Number of Comments	Decision
0	xn3aad	loseit	Why am I gaining weight on a healthy diet?	Hello, \nI really need some advice for someone...	Shinigami_Sadies	9/24/2022 16:33	9/24/2022 16:33	2022_9_23	29.0	0.0
1	xyz6l9	xxfitness	I ran a sub 30 5K	I know that's easy for most people but I haven...	Choice_Ad522	10/8/2022 13:59	10/8/2022 13:59	2022_10_7	3.0	1.0
2	x6u9cp	loseit	I am losing my mind instead of weight >:(I am currently 276, so 79 pounds down from my ...	peppermintwhitemocha	9/5/2022 19:09	9/5/2022 19:09	2022_9_4	22.0	0.0
3	ybse4o	bodybuilding	Please answer!!!	Do men bodybuilders use waist trainers?????	GrowthMobile	10/23/2022 16:41	10/23/2022 16:42	2022_10_22	12.0	1.0
4	x0yt2o	loseit	Worried about muscle loss	Hey\nSo basically I'm a 30 year old male Im 6f...	Takedownkd	8/29/2022 16:22	8/29/2022 16:22	2022_8_28	4.0	0.0

In []:

```
state = df.groupby('Subreddits')[['Number of Comments']].agg(['sum', 'mean', 'count']).reset_index()
sorted_state = state.sort_values(by=('Number of Comments', 'mean'), ascending=False)
print(sorted_state)
```

```
Subreddits Number of Comments
sum mean count
0 EatCheapAndHealthy 4909.0 24.918782 197
8 vegan 6430.0 19.484848 330
7 orangetheory 6551.0 18.610795 352
2 bodyweightfitness 5107.0 17.794425 287
5 nattyorjuice 546.0 14.368421 38
3 crossfit 2112.0 14.174497 149
6 nutrition 4710.0 12.975207 363
4 loseit 12326.0 11.595484 1063
9 xxfitness 4491.0 9.980000 450
1 bodybuilding 2194.0 9.838565 223
```

In []:

```
# Merge two columns
df['Titles_Body'] = df['Titles']+' '+df['Body']
df.head()
```

Out []:

	Reddit_ID	Subreddits	Titles	Body	Author	Initial Create	Retrieved_date	Date_of_collection	Number of Comments	Decision	Titles_Body
0	xn3aad	loseit	Why am I gaining weight on a healthy diet?	Hello, \nI really need some advice for someon...	Shinigami_Sadies	9/24/2022 16:33	9/24/2022 16:33	2022_9_23	29.0	0.0	Why am I gaining weight on a healthy diet? Hel...
1	xyz6l9	xxfitness	I ran a sub 30 5K	I know that's easy for most people but I haven...	Choice_Ad522	10/8/2022 13:59	10/8/2022 13:59	2022_10_7	3.0	1.0	I ran a sub 30 5K I know that's easy for most ...
2	x6u9cp	loseit	I am losing my mind instead of weight >:(I am currently 276, so 79 pounds down from my ...	peppermintwhitemocha	9/5/2022 19:09	9/5/2022 19:09	2022_9_4	22.0	0.0	I am losing my mind instead of weight >:(I...
3	ybse4o	bodybuilding	Please answer!!!	Do men bodybuilders use waist trainers?????	GrowthMobile	10/23/2022 16:41	10/23/2022 16:42	2022_10_22	12.0	1.0	Please answer!!! Do men bodybuilders use waist...
4	x0yt2o	loseit	Worried about muscle loss	Hey\nSo basically I'm a 30 year old male Im 6f...	Takedownkd	8/29/2022 16:22	8/29/2022 16:22	2022_8_28	4.0	0.0	Worried about muscle loss Hey\nSo basically I'...

```

In [ ]: # Customize your pre-processing
def preprocess_text(text):
    # remove Punctuation
    text = re.sub(r'^\w\s', '', text)
    # tokenize
    tokens = word_tokenize(text)
    # convert to lowercase
    tokens = [t.lower() for t in tokens]
    # remove stopwords
    stop_words = set(stopwords.words('english'))
    tokens = [t for t in tokens if t not in stop_words]
    # stemming
    stemmer = PorterStemmer()

```

```
tokens = [stemmer.stem(t) for t in tokens]
return tokens
```

```
In [ ]: # Pre-process post content
df['processed'] = df['Titles_Body'].apply(preprocess_text)
```

```
In [ ]: # Double check pre-processed posts
print(df['processed'][:10])
```

```
0    [gain, weight, healthi, diet, hello, realli, n...
1    [ran, sub, 30, 5k, know, that, easi, peopl, ha...
2    [lose, mind, instead, weight, gt, current, 276...
3    [pleas, answer, men, bodybuild, use, waist, tr...
4    [worri, muscl, loss, hey, basic, im, 30, year,...
5    [mount, pull, bar, miss, screw, recent, bought...
6    [5, week, alreadi, see, differ, small, mention...
7    [healthi, cocacola, make, skinni, fat, eat, lo...
8    [sick, cant, go, gym, sad, h, cope, happen]
9    [lose, weight, even, massiv, calori, deficit, ...
Name: processed, dtype: object
```

```
In [ ]: # Count Word frequency
word_counts = Counter()
for word in df['processed']:
    word_counts.update(word)
word_counts
```

```
Out[ ]: Counter({'gain': 778,  
                'weight': 3058,  
                'healthi': 528,  
                'diet': 786,  
                'hello': 178,  
                'realli': 1124,  
                'need': 930,  
                'advic': 533,  
                'someon': 268,  
                'pco': 25,  
                'loos': 138,  
                'last': 564,  
                'year': 1153,  
                'final': 188,  
                'reach': 161,  
                'goal': 587,  
                'ideal': 60,  
                'howev': 250,  
                'eat': 2316,  
                'littl': 456,  
                'food': 997,  
                'feint': 2,  
                'lot': 716,  
                'time': 1636,  
                'life': 434,  
                'end': 390,  
                'bmi': 79,  
                'doctor': 120,  
                'said': 214,  
                'goodbut': 1,  
                'knew': 53,  
                'wasnt': 150,  
                'decid': 207,  
                'better': 408,  
                'go': 1620,  
                'focus': 93,  
                'nutrit': 151,  
                'right': 466,  
                'type': 120,  
                'rather': 89,  
                'less': 274,  
                'even': 827,  
                'though': 293,  
                'good': 836,  
                'nolong': 1,  
                'rapidli': 9,  
                'notic': 219,  
                'peopl': 640,  
                'care': 136,  
                'due': 219,
```

'health': 301,
'risk': 29,
'warn': 17,
'activ': 262,
'8': 249,
'hour': 325,
'work': 1475,
'feet': 46,
'break': 134,
'5': 552,
'day': 2436,
'week': 1720,
'workday': 1,
'4am': 2,
'vegetablelentil': 1,
'soup': 49,
'1': 862,
'tin': 5,
'8am': 6,
'hand': 135,
'grape': 13,
'2': 898,
'rice': 163,
'cake': 49,
'low': 307,
'fat': 859,
'12': 337,
'small': 211,
'wholegrain': 4,
'bread': 112,
'roll': 43,
'beetroot': 3,
'salad': 78,
'tomato': 54,
'butter': 75,
'spread': 17,
'2pm': 2,
'eggcup': 1,
'size': 174,
'amount': 229,
'nut': 44,
'4pm': 1,
'stirfri': 4,
'veg': 32,
'bean': 67,
'teaspoon': 17,
'cook': 309,
'weekend': 121,
'usual': 324,
'dinner': 185,

'lunch': 153,
'vari': 30,
'breakfast': 118,
'toast': 18,
'sometim': 209,
'snack': 187,
'crisp': 7,
'drink': 309,
'tea': 52,
'coffe': 51,
'water': 334,
'help': 924,
'chang': 453,
'vegetarian': 95,
'dairi': 54,
'give': 353,
'acn': 14,
'tri': 1369,
'consum': 132,
'possibl': 235,
'keep': 496,
'like': 2257,
'fear': 31,
'effect': 115,
'mental': 152,
'poorli': 6,
'ran': 29,
'sub': 89,
'30': 602,
'5k': 15,
'know': 1360,
'that': 378,
'easi': 158,
'havent': 235,
'abl': 314,
'run': 480,
'56': 69,
'lose': 1380,
'mind': 152,
'instead': 152,
'gt': 64,
'current': 401,
'276': 4,
'79': 3,
'pound': 517,
'highest': 38,
'355': 1,
'super': 168,
'proud': 82,
'omg': 4,

'month': 967,
'fix': 88,
'around': 657,
'1500': 79,
'calori': 1440,
'use': 1030,
'scale': 260,
'mfp': 20,
'track': 321,
'mile': 214,
'walk': 449,
'train': 512,
'3': 795,
'45': 279,
'min': 541,
'ellipt': 27,
'trail': 9,
'saturday': 64,
'zumba': 4,
'sunday': 59,
'ive': 2078,
'gotten': 83,
'desper': 28,
'ask': 320,
'adipex': 3,
'see': 590,
'would': 1316,
'doesnt': 246,
'im': 4447,
'sure': 384,
'next': 272,
'got': 528,
'ta': 14,
'pleas': 294,
'answer': 141,
'men': 47,
'bodybuild': 53,
'waist': 54,
'trainer': 72,
'worri': 136,
'muscl': 696,
'loss': 500,
'hey': 176,
'basic': 200,
'old': 250,
'male': 94,
'6ft1': 2,
'weigh': 358,
'88kg': 4,
'10': 422,

'ago': 383,
'gym': 789,
'4': 444,
'90130': 1,
'g': 17,
'protein': 612,
'depend': 58,
'manag': 122,
'come': 331,
'78kg': 4,
'bodi': 1001,
'percentag': 40,
'1918': 1,
'percent': 10,
'accord': 55,
'boditrax': 1,
'machin': 123,
'accur': 63,
'follow': 217,
'trend': 17,
'show': 155,
'lost': 606,
'4kg': 5,
'sinc': 629,
'start': 1587,
'free': 124,
'mass': 97,
'drop': 143,
'3kg': 4,
'much': 1021,
'normal': 266,
'workout': 865,
'increas': 223,
'massiv': 26,
'deadlift': 104,
'40kg': 5,
'80kg': 11,
'calv': 36,
'rais': 145,
'85': 24,
'kg': 209,
'exampl': 104,
'shoukd': 1,
'trust': 38,
'seem': 383,
'consist': 234,
'hasnt': 41,
'gone': 94,
'appreci': 256,
'overthink': 11,

```

'let': 256,
'thank': 564,
'mount': 4,
'pull': 177,
'bar': 143,
'miss': 110,
'screw': 16,
'recent': 309,
'bought': 72,
'friend': 284,
'lead': 53,
'avoid': 116,
'death': 25,
'mean': 208,
'there': 157,
'fine': 115,
'barhttpswwwamazoncoukonetwofitmountedstableoutdoormaximumdpb07v66qq37_encodingutf8amppd_rd_whzrvdampcontentidamzn1sym31c9090c9b654f91bf3704dd92281172amppf_rd_p31c9090c9b654f91bf3704dd92281172amppf_rd_rq84s95zrx8pnwjzapzyhamppd_rd_wgasssvamppd_rd_r4c71ac1c2c1d486d8878a4aad9e61dcdampref_pd_gw_ci_mcx_mihttpswwwamazoncoukonetwofitmountedstableoutdoormaximumdpb07v66qq37_encodingutf8amppd_rd_whzrvdampcontentidamzn1sym31c9090c9b654f91bf3704dd92281172amppf_rd_p31c9090c9b654f91bf3704dd92281172amppf_rd_rq84s95zrx8pnwjzapzyhamppd_r4c71ac1c2c1d486d8878a4aad9e61dcdampref_pd_gw_ci_mcx_mi': 1,
'alreadi': 208,
'differ': 361,
'mention': 73,
'ed': 21,
'sv': 13,
'18intersex': 1,
'first': 585,
'118': 3,
'260': 11,
'lb': 519,
'sad': 45,
'anyth': 368,
'motiv': 203,
'journey': 230,
'113': 1,
'249': 2,
'remark': 1,
'easier': 78,
'past': 327,
'attempt': 61,
'pleasantli': 3,
'surpris': 41,
'get': 2008,
'hate': 152,
'could': 527,
'form': 151,
'coher': 1,
'thought': 375,
'turn': 135,

```

'healthiest': 15,
'way': 666,
'easiest': 7,
'albeit': 3,
'slowest': 4,
'starv': 46,
'think': 901,
'look': 1010,
'envi': 3,
'still': 759,
'struggl': 224,
'one': 954,
'tbh': 18,
'far': 188,
'face': 107,
'defin': 19,
'okay': 85,
'definit': 148,
'genet': 52,
'bone': 63,
'structur': 32,
'thing': 647,
'cheekbon': 2,
'pronounc': 5,
'also': 990,
'favorit': 80,
'pair': 47,
'jean': 38,
'looser': 6,
'thigh': 63,
'great': 252,
'fit': 506,
'kind': 201,
'word': 89,
'post': 406,
'posit': 82,
'benefici': 10,
'wish': 56,
'everyon': 285,
'luck': 19,
'cocacola': 2,
'make': 1031,
'skinni': 146,
'junk': 79,
'liter': 89,
'cola': 2,
'185': 19,
'cm': 32,
'60': 132,
'quit': 243,

'went': 347,
'70': 38,
'want': 1591,
'sick': 84,
'cant': 579,
'h': 12,
'cope': 15,
'happen': 210,
'deficit': 277,
'exercis': 875,
'regularli': 50,
'stupid': 34,
'homework': 3,
'lol': 118,
'higher': 70,
'intak': 160,
'lower': 227,
'least': 208,
'dont': 1690,
'immedi': 51,
'back': 1080,
'minut': 588,
'1000': 39,
'shit': 68,
'metabol': 89,
'what': 121,
'two': 363,
'mum': 13,
'put': 296,
'oil': 115,
'soemth': 1,
'lmfao': 1,
'anyon': 712,
'els': 275,
'experienc': 57,
'rebuild': 4,
'routin': 401,
'lss': 1,
'long': 452,
'stori': 92,
'short': 132,
'pretti': 379,
'crossfit': 96,
'covid': 78,
'5x': 17,
'2x': 20,
'sprain': 5,
'ankl': 41,
'pt': 32,
'releas': 21,

'juli': 39,
'find': 462,
'demotiv': 12,
'stress': 127,
'balanc': 81,
'strategi': 52,
'found': 173,
'build': 256,
'ever': 241,
'made': 330,
'video': 100,
'cbum': 1,
'say': 487,
'btch': 1,
'connect': 17,
'hrm': 4,
'non': 42,
'tablet': 13,
'rower': 68,
'occasion': 76,
'visit': 32,
'studio': 60,
'home': 322,
'otf': 183,
'1kg': 9,
'2lb': 9,
'hi': 315,
'13yearold': 1,
'unabl': 20,
'felt': 232,
'imposs': 37,
'grandma': 10,
'almost': 226,
'crazi': 76,
'portion': 103,
'person': 258,
'smaller': 57,
'scold': 1,
'enjoy': 163,
'take': 642,
'tell': 249,
'ye': 55,
'asian': 13,
'insan': 31,
'per': 335,
'58': 31,
'larg': 97,
'bowl': 51,
'hard': 388,
'convinc': 33,

'parent': 64,
'buy': 205,
'ton': 67,
'brother': 32,
'told': 152,
'number': 149,
'benefit': 56,
'famili': 162,
'overweight': 134,
'physic': 148,
'everyday': 127,
'enough': 344,
'fight': 37,
'urg': 16,
'boredom': 12,
'overeate': 16,
'tast': 90,
'someth': 535,
'ill': 332,
'control': 136,
'feel': 1684,
'unsatisfi': 3,
'frustrat': 112,
'lack': 67,
'progress': 388,
'shape': 103,
'advanc': 111,
'tip': 234,
'jog': 30,
'speed': 50,
'orangetheori': 20,
'runner': 32,
'previous': 33,
'steadili': 18,
'1012': 9,
'base': 522,
'pace': 85,
'inclin': 111,
'push': 609,
'ao': 239,
'templat': 81,
'maintain': 162,
'overal': 81,
'especi': 171,
'improv': 155,
'everi': 462,
'class': 341,
'arent': 89,
'quad': 32,
'leg': 401,

'52': 37,
'mediumshort': 1,
'height': 90,
'wonder': 288,
'attract': 22,
'proport': 11,
'thoughtsexperi': 2,
'chocol': 62,
'mani': 348,
'slightli': 66,
'silli': 15,
'question': 457,
'approxim': 19,
'550': 2,
'100': 130,
'gram': 69,
'matter': 131,
'dark': 17,
'milk': 167,
'white': 59,
'read': 198,
'ingredi': 89,
'cocoa': 6,
'228': 3,
'cal': 126,
'sugar': 189,
'387': 1,
'42': 6,
'884': 1,
'idea': 287,
'everyth': 243,
'except': 80,
'process': 96,
'product': 176,
'individu': 18,
'imagin': 40,
'option': 197,
'five': 35,
'yet': 147,
'53': 31,
'21': 53,
'f': 43,
'180': 35,
'count': 261,
'loseit': 18,
'app': 170,
'stay': 294,
'treadmil': 118,
'46': 17,
'wan': 91,

'na': 158,
'165': 20,
'anniversari': 3,
'mid': 45,
'novemb': 19,
'gon': 66,
'happi': 184,
'done': 228,
'83': 3,
'updat': 47,
'alway': 477,
'obes': 80,
'anorex': 5,
'age': 80,
'17': 38,
'1101b': 7,
'59': 43,
'bdd': 1,
'slim': 44,
'18': 65,
'20': 225,
'1351451b': 1,
'stop': 368,
'restrict': 75,
'ate': 150,
'fast': 299,
'bit': 427,
'hit': 225,
'symptom': 14,
'certain': 65,
'issu': 272,
'battl': 12,
'psychosi': 1,
'physiolog': 2,
'1021b': 1,
'heaviest': 29,
'2371b': 3,
'antipsychot': 1,
'didnt': 352,
'either': 172,
'anyway': 135,
'june': 44,
'particular': 26,
'incid': 6,
'click': 46,
'realiz': 108,
'never': 459,
'spectrum': 3,
'opt': 6,
'approach': 55,

'join': 121,
'noom': 13,
'sponsor': 2,
'log': 67,
'emot': 50,
'251b': 5,
'middl': 25,
'begin': 100,
'septemb': 89,
'camp': 7,
'trip': 31,
'anti': 8,
'depress': 121,
'101b': 24,
'anymor': 123,
'1181b': 3,
'plan': 316,
'point': 284,
'dress': 27,
'mom': 67,
'wed': 36,
'summer': 55,
'tldr': 36,
'anorexia': 12,
'coupl': 203,
'1011b': 1,
'191b': 2,
'1351b': 7,
'squat': 381,
'lighthead': 5,
'stick': 151,
'press': 201,
'experi': 192,
'close': 129,
'bedtim': 4,
'254': 2,
'caus': 195,
'addit': 59,
'myth': 11,
'700': 22,
'17502000': 1,
'late': 123,
'shower': 16,
'mainten': 120,
'300': 24,
'rich': 19,
'bag': 77,
'dri': 67,
'pineappl': 4,
'bed': 72,

'replac': 63,
'oat': 52,
'chicken': 247,
'peanut': 31,
'safest': 2,
'absolut': 81,
'total': 204,
'eliminationa': 1,
'carbveggiefruit': 1,
'sourc': 64,
'choos': 64,
'categori': 17,
'refreez': 4,
'begond': 1,
'brick': 6,
'stock': 17,
'beyond': 32,
'beef': 56,
'refridger': 1,
'section': 35,
'sale': 17,
'bec': 2,
'bacteria': 3,
'degrad': 5,
'true': 61,
'safe': 56,
'sauc': 72,
'amp': 265,
'chilli': 5,
'patti': 3,
'textur': 23,
'might': 193,
'ok': 84,
'bunch': 49,
'freez': 33,
'strength': 323,
'without': 302,
'mess': 50,
'id': 370,
'pullup': 86,
'weaken': 2,
'arm': 235,
'up': 208,
'may': 218,
'averag': 136,
'well': 458,
'compar': 78,
'refer': 45,
'assist': 27,
'focu': 145,

'em': 8,
'weaker': 17,
'aspect': 21,
'full': 215,
'skill': 41,
'practic': 65,
'access': 48,
'hero': 8,
'wod': 45,
'typic': 58,
'tough': 61,
'modifi': 16,
'shorter': 24,
'side': 190,
'tg': 1,
'candid': 8,
'ear': 13,
'recommend': 260,
'vein': 14,
'fell': 32,
'left': 148,
'exterior': 1,
'calf': 44,
'bottom': 25,
'38': 13,
'titl': 67,
'realis': 26,
'1901b': 10,
'big': 263,
'set': 404,
'mileston': 30,
'along': 63,
'today': 563,
'reddit': 48,
'alon': 52,
'share': 342,
'success': 83,
'interest': 127,
'cut': 292,
'empti': 33,
'live': 244,
'kept': 55,
'aim': 54,
'1000day': 1,
'2201b': 5,
'finger': 31,
'cross': 31,
'bing': 158,
'tini': 27,
'vacat': 39,

'3k': 4,
'strenght': 10,
'ham': 15,
'macro': 105,
'23': 90,
'kilo': 27,
'subtl': 1,
'measur': 95,
'hopeless': 20,
'miser': 27,
'ruin': 39,
'continu': 161,
'recomp': 34,
'mayb': 315,
'intuit': 14,
'religi': 17,
'pal': 11,
'platform': 6,
'samsung': 8,
'goe': 70,
'target': 46,
'near': 53,
'1850150': 1,
'1633': 2,
'add': 219,
'whatev': 113,
'fitnessexercis': 2,
'earn': 7,
'lift': 511,
'40': 101,
'600': 20,
'2233': 1,
'wherea': 8,
'1850': 2,
'm21': 1,
'163': 3,
'impost': 1,
'syndrom': 5,
'mirror': 43,
'isnt': 215,
'weird': 85,
'thursday': 57,
'submit': 3,
'check': 165,
'program': 239,
'saw': 71,
'growth': 34,
'anadrom': 1,
'50': 102,
'medic': 93,

'buydu': 1,
'legit': 8,
'new': 422,
'world': 73,
'curiou': 87,
'meet': 38,
'man': 58,
'6': 439,
'ft': 13,
'195': 9,
'10000': 17,
'step': 197,
'doubl': 28,
'prework': 2,
'women': 77,
'taken': 57,
'tdee': 99,
'bmrrmr': 1,
'onlin': 99,
'calcul': 105,
'regard': 45,
'burn': 172,
'resourc': 37,
'best': 285,
'estim': 29,
'throughout': 78,
'harder': 68,
'actual': 331,
'complet': 212,
'odd': 20,
'round': 174,
'sedentari': 61,
'bmr': 22,
'cool': 48,
'lay': 31,
'coma': 1,
'clean': 110,
'movement': 92,
'etc': 308,
'neat': 5,
'thermic': 1,
'200300': 4,
'wat': 1,
'account': 276,
'top': 111,
'15000': 3,
'17000': 1,
'reliabl': 8,
'seemingli': 6,
'credibl': 1,

'watch': 136,
'strap': 15,
'wrist': 42,
'abund': 2,
'octob': 72,
'benchmark': 26,
'throw': 59,
'anxieti': 55,
'stomach': 138,
'anxiou': 19,
'limit': 126,
'rariti': 1,
'regain': 15,
'scare': 72,
'high': 349,
'daili': 459,
'threw': 15,
'pizza': 51,
'whole': 209,
'asid': 27,
'guilti': 28,
'30min': 10,
'cocain': 1,
'overwhelm': 25,
'panic': 11,
'eaten': 49,
'suddenli': 27,
'rid': 36,
'toilet': 6,
'coke': 8,
'guess': 146,
'roast': 43,
'carb': 202,
'panik': 1,
'randomli': 18,
'bathroom': 10,
'wouldnt': 83,
'tomorrow': 68,
'bad': 342,
'garlic': 39,
'wrap': 17,
'sensat': 6,
'fatter': 10,
'alley': 1,
'morbidli': 8,
'pressur': 32,
'poor': 32,
'excus': 27,
'unemploy': 4,
'gf': 10,


'anywher': 29,
'hell': 87,
'wrong': 183,
'anxietydrinkingdrug': 1,
'usetoo': 1,
'purpos': 25,
'bulimia': 5,
'ref': 2,
'26': 46,
'62': 24,
'1941b': 4,
'330': 7,
'unpleas': 1,
'smell': 29,
'attend': 29,
'member': 93,
'becam': 40,
'frequent': 41,
'particip': 20,
'strong': 83,
'odor': 1,
'wors': 85,
'front': 101,
'desk': 23,
'remind': 27,
'facebook': 5,
'instagram': 8,
'sign': 49,
'shrink': 15,
'seen': 120,
'board': 12,
'oh': 41,
'neg': 45,
'tornado': 22,
'becom': 155,
'deal': 94,
'unwelcom': 1,
'deserv': 11,
'shake': 73,
'appl': 61,
'meal': 590,
'troubl': 59,
'planningrememb': 1,
'eatfeel': 1,
'regular': 66,
'forc': 65,
'chug': 2,
'caffein': 24,
'honest': 33,
'questionpleas': 1,

'asshol': 3,
'vegan': 889,
'lifestyl': 100,
'relat': 67,
'truli': 34,
'impact': 50,
'wont': 160,
'anyjust': 1,
'genuin': 44,
'disord': 68,
'16ftm': 2,
'169cm': 2,
'sw': 53,
'68': 23,
'cw': 56,
'49': 4,
'cheat': 75,
'half': 204,
'750': 2,
'monday': 70,
'400': 16,
'uncomfort': 43,
'reason': 212,
'unreli': 2,
'nondigit': 1,
'15': 262,
'horror': 3,
'weighin': 69,
'school': 142,
'nurs': 6,
'499': 1,
'pant': 40,
'cri': 33,
'expect': 76,
'tmi': 4,
'poop': 7,
'serious': 66,
'hope': 266,
'underweight': 33,
'mostli': 151,
'allergi': 10,
'condit': 42,
'veggi': 121,
'tire': 149,
'carrot': 26,
'celeri': 7,
'cucumb': 17,
'snap': 15,
'pea': 22,
'crunch': 26,

```
'kinda': 78,  
'brain': 45,  
'chew': 19,  
'tasti': 20,  
...})
```

```
In [ ]: # Top 20 words  
most_common_words = word_counts.most_common(20)  
words, frequencies = zip(*most_common_words)  
print(words)  
print(frequencies)  
  
(  
'im', 'weight', 'day', 'eat', 'like', 'ive', 'get', 'week', 'dont', 'feel', 'time', 'go', 'want', 'start', 'work', 'calori', 'lose',  
'tri', 'know', 'would')  
(4447, 3058, 2436, 2316, 2257, 2078, 2008, 1720, 1690, 1684, 1636, 1620, 1591, 1587, 1475, 1440, 1380, 1369, 1360, 1316)
```

```
In [ ]: # Word frequency visualization  
# Word frequency visualization  
plt.figure(figsize = (10, 6))  
plt.barh(words, frequencies)  
plt.xlabel('Frequencies')  
plt.ylabel('Words')  
plt.title('Top 20 Most Common Words')  
plt.gca().invert_yaxis()  
plt.show()
```



```
In [ ]: # Word Cloud
all_words = [word for list in df['processed'] for word in list]
text = ' '.join(all_words)
wordcloud = WordCloud(max_words = 200, background_color = 'white').generate(text)
plt.figure(figsize = (10,5))
plt.imshow(wordcloud)
plt.show()
```



In []: text

```
In [ ]: #create a dictionary
dictionary = corpora.Dictionary(df['processed'])
#create a corpus
corpus = [dictionary.doc2bow(text) for text in df['processed']]
```

```
In [ ]: # train LDA, Topic = 2
lda_model_2 = LdaModel(corpus = corpus, id2word = dictionary, num_topics = 2, random_state = 42)
# topic word matrix
for idx, topic in lda_model_2.print_topics (-1, num_words = 10):
    print('Topic: {} \nwords: {}'.format(idx, topic))
```

```
WARNING:gensim.models.ldamodel:too few updates, training might not converge; consider increasing the number of passes or iterations to improve accuracy
```

```
Topic: 0
words: 0.017*"im" + 0.013*"weight" + 0.012*"day" + 0.009*"eat" + 0.008*"get" + 0.008*"like" + 0.007*"week" + 0.007*"feel" + 0.007*"start" + 0.007*"calori"
Topic: 1
words: 0.014*"im" + 0.010*"ive" + 0.008*"weight" + 0.008*"like" + 0.008*"eat" + 0.006*"time" + 0.006*"work" + 0.006*"go" + 0.006*"don't" + 0.006*"get"
```

```
In [ ]: # train LDA, Topic = 4
lda_model_4 = LdaModel(corpus = corpus, id2word = dictionary, num_topics = 4, random_state = 42)
# topic word matrix
for idx, topic in lda_model_4.print_topics (-1, num_words = 10):
    print('Topic: {} \nwords: {}'.format(idx, topic))
```

WARNING:gensim.models.ldamodel:too few updates, training might not converge; consider increasing the number of passes or iterations to improve accuracy

Topic: 0

words: 0.018*"im" + 0.013*"weight" + 0.013*"day" + 0.009*"get" + 0.008*"like" + 0.007*"start" + 0.007*"feel" + 0.006*"dont" + 0.006*"f
at" + 0.006*"eat"

Topic: 1

words: 0.014*"im" + 0.008*"like" + 0.008*"weight" + 0.006*"ive" + 0.006*"go" + 0.006*"get" + 0.006*"time" + 0.006*"dont" + 0.006*"bod
i" + 0.006*"tri"

Topic: 2

words: 0.011*"im" + 0.008*"1" + 0.007*"like" + 0.007*"min" + 0.007*"sec" + 0.007*"time" + 0.007*"2" + 0.006*"row" + 0.006*"3" + 0.006
*"week"

Topic: 3

words: 0.019*"im" + 0.015*"weight" + 0.014*"eat" + 0.012*"ive" + 0.010*"day" + 0.009*"want" + 0.008*"like" + 0.008*"lose" + 0.008*"wee
k" + 0.008*"calori"

```
In [ ]: # train LDA, Topic = 8
lda_model_8 = LdaModel(corpus = corpus, id2word = dictionary, num_topics = 8, random_state = 42)
# topic word matrix
for idx, topic in lda_model_8.print_topics (-1, num_words = 10):
    print('Topic: {} \nwords: {}'.format(idx, topic))
```

WARNING:gensim.models.ldamodel:too few updates, training might not converge; consider increasing the number of passes or iterations to improve accuracy

Topic: 0
words: 0.012*"im" + 0.010*"day" + 0.008*"question" + 0.007*"like" + 0.007*"get" + 0.007*"account" + 0.006*"share" + 0.006*"daili" + 0.005*"dont" + 0.005*"thread"

Topic: 1
words: 0.011*"im" + 0.010*"vegan" + 0.008*"time" + 0.007*"like" + 0.006*"anyon" + 0.006*"go" + 0.005*"tri" + 0.004*"get" + 0.004*"work" + 0.004*"need"

Topic: 2
words: 0.015*"sec" + 0.015*"1" + 0.014*"min" + 0.013*"block" + 0.013*"x" + 0.013*"row" + 0.011*"base" + 0.011*"push" + 0.010*"30" + 0.010*"minut"

Topic: 3
words: 0.015*"vegan" + 0.014*"im" + 0.010*"eat" + 0.009*"work" + 0.009*"want" + 0.008*"like" + 0.007*"ive" + 0.007*"food" + 0.006*"day" + 0.006*"anim"

Topic: 4
words: 0.018*"im" + 0.016*"eat" + 0.014*"like" + 0.011*"dont" + 0.010*"feel" + 0.009*"calori" + 0.009*"protein" + 0.008*"get" + 0.007*"day" + 0.007*"fat"

Topic: 5
words: 0.019*"weight" + 0.014*"im" + 0.012*"week" + 0.012*"ive" + 0.010*"eat" + 0.010*"start" + 0.009*"day" + 0.009*"lose" + 0.009*"get" + 0.009*"like"

Topic: 6
words: 0.022*"im" + 0.013*"weight" + 0.010*"eat" + 0.010*"day" + 0.007*"calori" + 0.007*"ive" + 0.007*"diet" + 0.007*"feel" + 0.006*"tri" + 0.006*"get"

Topic: 7
words: 0.022*"im" + 0.015*"weight" + 0.012*"day" + 0.009*"dont" + 0.008*"ive" + 0.008*"get" + 0.008*"like" + 0.008*"feel" + 0.007*"start" + 0.007*"eat"

```
In [ ]: # LDA evaluation
# Coherence
coherence_mode11_lda = CoherenceModel (model = lda_model_2, texts = df['processed'],
                                         dictionary = dictionary, coherence = 'c_v').get_coherence()

coherence_mode12_lda = CoherenceModel (model = lda_model_4, texts = df['processed'],
                                         dictionary = dictionary, coherence = 'c_v').get_coherence()

coherence_mode13_lda = CoherenceModel (model = lda_model_8, texts = df['processed'],
                                         dictionary = dictionary, coherence = 'c_v').get_coherence()

print('\nCoherence score 1', coherence_mode11_lda)
print('\nCoherence score 2', coherence_mode12_lda)
print('\nCoherence score 3', coherence_mode13_lda)
```

Coherence score 1 0.5352799718462801
InCoherence score 2 0.4897607824133229

Coherence score 3 0.5352799718462801

Sentiment Analysis

```
In [ ]: # Define sentiment analysis function
analyzer = SentimentIntensityAnalyzer()
```

```
def get_post_sentiment(text):
    return analyzer.polarity_scores(text)['compound'] # return sentiment of a tweet
df['sentiment_score_vader'] = df['Titles_Body'].apply(get_post_sentiment)
df.head()
```


Out []:

	Reddit_ID	Subreddits	Titles	Body	Author	Initial Create	Retrieved_date	Date_of_collection	Number of Comments	Decision	Titles_Body	pro
0	xn3aad	loseit	Why am I gaining weight on a healthy diet?	Hello, \nI really need some advice for someone...	Shinigami_Sadies	9/24/2022 16:33	9/24/2022 16:33	2022_9_23	29.0	0.0	Why am I gaining weight on a healthy diet? Hel...	die re
1	xyz6l9	xxfitness	I ran a sub 30 5K	I know that's easy for most people but I haven...	Choice_Ad522	10/8/2022 13:59	10/8/2022 13:59	2022_10_7	3.0	1.0	I ran a sub 30 5K I know that's easy for most ...	[r] th
2	x6u9cp	loseit	I am losing my mind instead of weight >:(I am currently 276, so 79 pounds down from my ...	peppermintwhitemocha	9/5/2022 19:09	9/5/2022 19:09	2022_9_4	22.0	0.0	I am losing my mind instead of weight >:(I...	i wei
3	ybse4o	bodybuilding	Please answer!!!	Do men bodybuilders use waist trainers?????	GrowthMobile	10/23/2022 16:41	10/23/2022 16:42	2022_10_22	12.0	1.0	Please answer!!! Do men bodybuilders use waist...	boc use
4	x0yt2o	loseit	Worried about muscle loss	Hey\nSo basically I'm a 30 year old male Im 6f...	Takedownkd	8/29/2022 16:22	8/29/2022 16:22	2022_8_28	4.0	0.0	Worried about muscle loss Hey\nSo basically I'...	lo b: 30

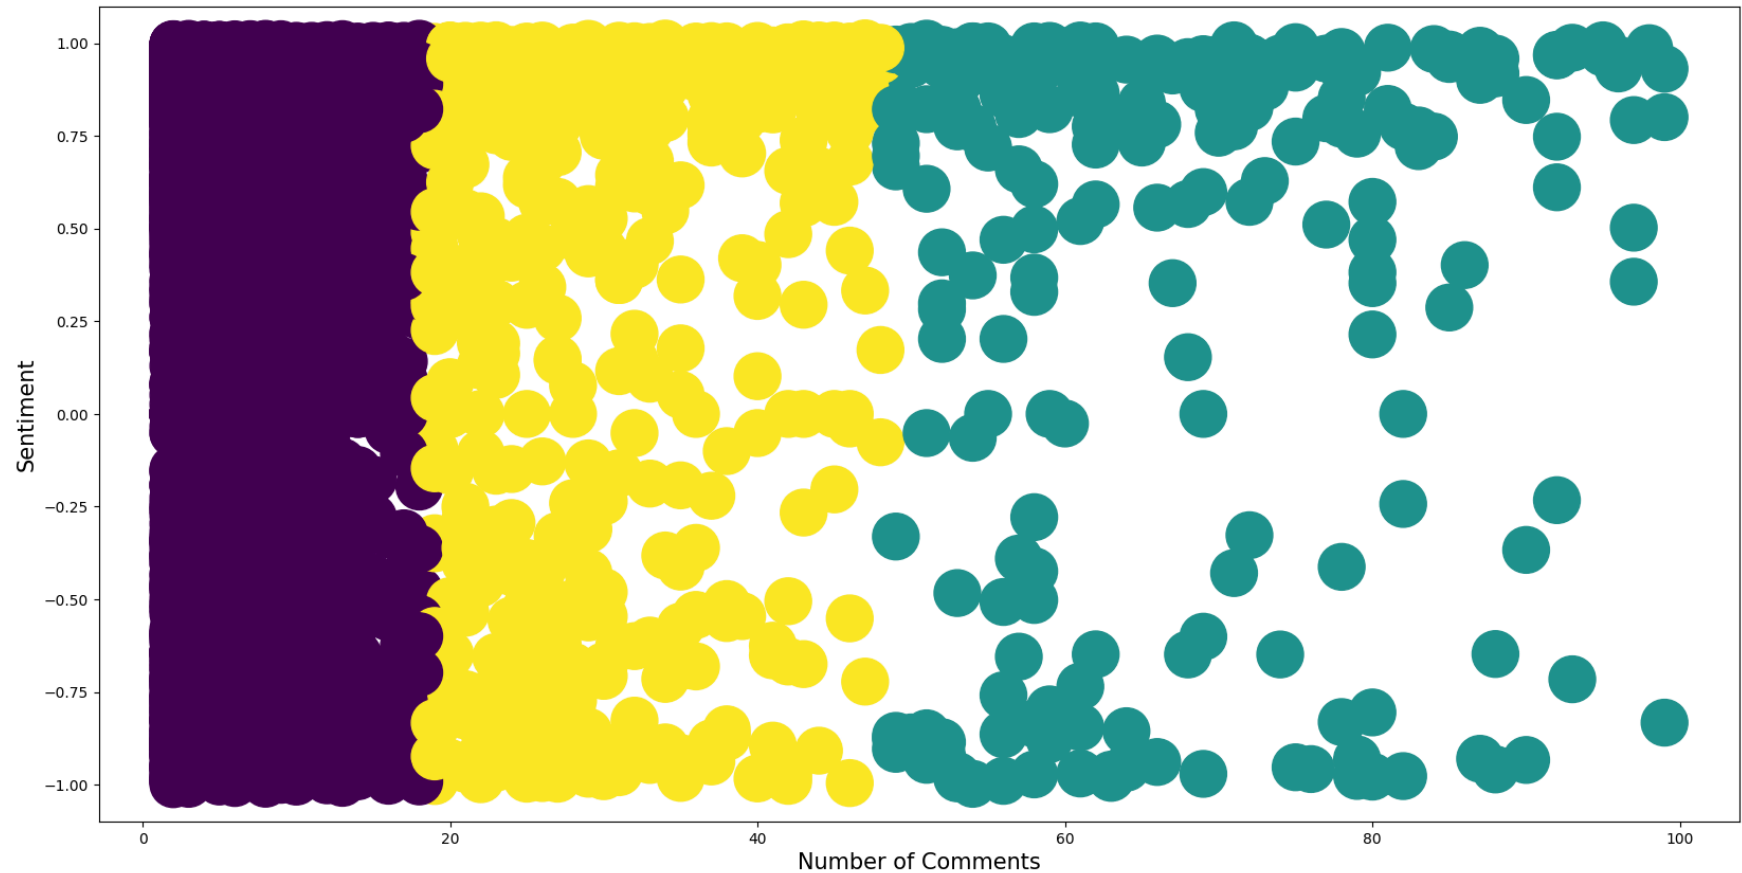
In []: #K-mean Clustering

In []: data = list(zip(df['Number of Comments'], df['sentiment_score_vader']))

```
inertias = []  
for i in range (1,11):  
    kmeans = KMeans(n_clusters=i)  
    kmeans.fit(data)  
    inertias.append (kmeans.inertia_)  
  
plt.figure(figsize=(12,6))  
plt.plot(range (1,11), inertias, marker='o')  
plt.xticks(np.arange(1, 11, 1.0),fontsize = 10)  
plt.yticks(fontsize = 10)  
plt.title('Elbow method',fontsize = 20)  
plt.xlabel('Number of clusters',fontsize = 15)  
plt.ylabel('WCSS',fontsize = 15)  
plt.show()
```




```
In [ ]: kmeans = KMeans(n_clusters=3)
kmeans.fit(data)
plt.figure(figsize=(20,10))
plt.scatter(df['Number of Comments'], df['sentiment_score_vader'], c=kmeans.labels_, s=1000)
plt.xlabel('Number of Comments', fontsize = 15)
plt.ylabel('Sentiment', fontsize = 15)
plt.xticks(fontsize = 10)
plt.yticks(fontsize = 10)
plt.show()
```



Hierarchical clustering

```
In [ ]: distance_matrix = pdist(df[['Number of Comments', 'sentiment_score_vader']], metric = "euclidean")
distance_matrix
```

```
Out[ ]: array([26.01016203,  7.01659556, 17.00331817, ..., 31.01471275,
        1.17919464, 32.00170461])
```

```
In [ ]: H_single = linkage(distance_matrix, method = 'single')
H_complete = linkage(distance_matrix, method = 'complete')
H_average = linkage(distance_matrix, method = 'average')
```

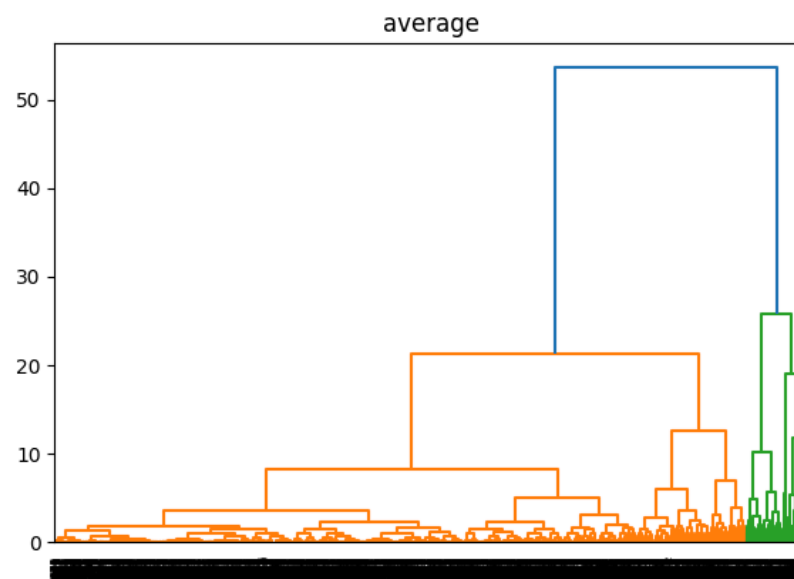


```
In [ ]: plt.figure(figsize = (15,10))

plt.subplot(2,2,1)
dendrogram(H_single)
plt.title("single")

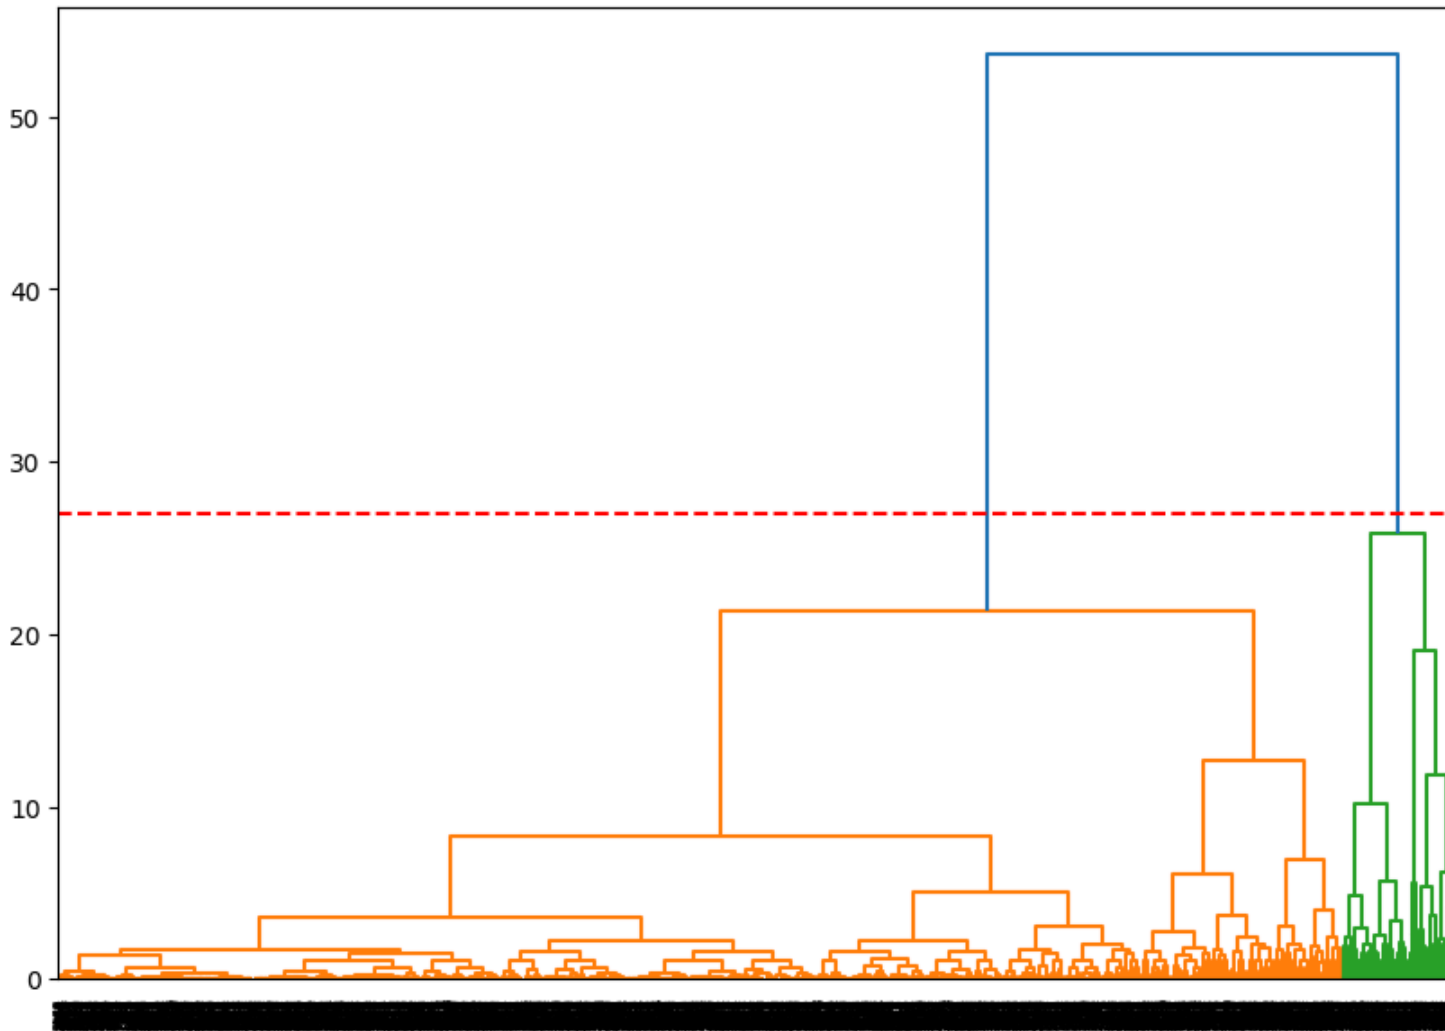
plt.subplot(2,2,2)
dendrogram(H_complete)
plt.title ("complete")

plt.subplot(2,2,3)
dendrogram(H_average)
plt.title("average")

plt.show()
```



```
In [ ]: # cut the dendrogram
groups = fcluster(H_average, t = 27, criterion = "maxclust")
plt.figure(figsize = (10,7))
dendrogram(H_average)
plt.axhline(y = 27, color = 'r', linestyle = "--")
plt.show()
```



Density-based Clustering

```
In [ ]: df[['Number of Comments', 'sentiment_score_vader']].shape
```

```
Out[ ]: (3452, 2)
```

```
In [ ]: db =DBSCAN(eps = 0.1,min_samples =3).fit(df[['Number of Comments','sentiment_score_vader']])
```

```
In [ ]: # Evaluation
sil_score_kmeans = silhouette_score(df[['Number of Comments', 'sentiment_score_vader']], kmeans.labels_)
sil_score_hc = silhouette_score(df[['Number of Comments', 'sentiment_score_vader']], groups)
sil_score_dbscan = silhouette_score(df[['Number of Comments', 'sentiment_score_vader']], db.labels_)
```

```
print(sil_score_kmeans)
print(sil_score_hc)
print(sil_score_dbscan)
```

```
0.6993406591002103
0.476129360170812
0.42662650847816064
```

Classification

```
In [ ]: # Define Predictors and Target variable
```

```
In [ ]: from sklearn.feature_extraction.text import TfidfVectorizer
import pandas as pd
import numpy as np

# Initialize TfidfVectorizer
tfidf_vectorizer = TfidfVectorizer(min_df=10, max_df=0.90)

# Convert phrases into lower case and apply TF-IDF vectorization
tfidf_dtm = tfidf_vectorizer.fit_transform(df['processed'].apply(lambda x: ' '.join(x) if isinstance(x, list) else x))

# Calculate sum of TF-IDF scores and get the indices of the top 40 terms
tfidf_sum_scores = tfidf_dtm.sum(axis=0).A1
tfidf_top_40_indices = tfidf_sum_scores.argsort()[-40:][::-1]

# Create the dataframe using only the columns for the top 40 terms
df_top_40_terms_tfidf = pd.DataFrame(tfidf_dtm[:, tfidf_top_40_indices].toarray(), columns=np.array(tfidf_vectorizer.get_feature_names_
df_top_40_terms_tfidf
```

Out []:

	im	weight	eat	day	ive	like	week	get	feel	calori	...	back	month	make	help	anyon	g
0	0.000000	0.226704	0.208797	0.100705	0.000000	0.031046	0.036923	0.000000	0.000000	0.000000	...	0.000000	0.000000	0.0	0.041935	0.0	0.0000
1	0.000000	0.000000	0.000000	0.261549	0.000000	0.000000	0.143844	0.000000	0.000000	0.000000	...	0.000000	0.000000	0.0	0.000000	0.0	0.0000
2	0.049363	0.244339	0.065637	0.253256	0.062922	0.000000	0.208924	0.000000	0.000000	0.079543	...	0.000000	0.080181	0.0	0.079096	0.0	0.0000
3	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.000000	0.0	0.000000	0.0	0.0000
4	0.094548	0.077999	0.125717	0.040423	0.080346	0.000000	0.177850	0.000000	0.000000	0.050785	...	0.000000	0.000000	0.0	0.000000	0.0	0.1136
...
3447	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.117483	0.000000	...	0.000000	0.000000	0.0	0.000000	0.0	0.0000
3448	0.000000	0.000000	0.088084	0.000000	0.084442	0.000000	0.093459	0.000000	0.000000	0.000000	...	0.109958	0.000000	0.0	0.106146	0.0	0.0000
3449	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.277998	0.000000	...	0.000000	0.000000	0.0	0.000000	0.0	0.1792
3450	0.000000	0.000000	0.070858	0.068351	0.000000	0.000000	0.075182	0.197518	0.000000	0.000000	...	0.000000	0.000000	0.0	0.000000	0.0	0.0000
3451	0.000000	0.097126	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.000000	0.0	0.000000	0.0	0.0000

3452 rows × 40 columns

```
In [ ]: X = df_top_40_terms_tfidf
        Y = df['Decision']
```

Data Partition

```
In [ ]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state = 42)
```

Decision Tree with pruning

```
In [ ]: from sklearn.tree import DecisionTreeClassifier

        # Initialize the DecisionTreeClassifier
        clf = DecisionTreeClassifier(random_state=42)

        # Train the classifier
        clf.fit(X_train, Y_train)
```

```
Out [ ]: DecisionTreeClassifier
        DecisionTreeClassifier(random_state=42)
```

```
In [ ]: # Fine-tune cost-complexity alpha value
path = clf.cost_complexity_pruning_path(X_train, Y_train)
ccp_alphas = path.ccp_alphas[:-1] # Exclude the maximum value; without any split
```

```
In [ ]: # create a table containing both error rate and ccp_alpha
# Obtain train and test scores using validation_curve
train_scores, test_scores = validation_curve (
    clf, X_train, Y_train, param_name="ccp_alpha", param_range=ccp_alphas, cv=5, scoring="accuracy"
)
train_error_rates = 1 - train_scores
test_error_rates = 1 - test_scores
# Display the results using ValidationCurveDisplay
display = ValidationCurveDisplay(
    param_name= 'ccp_alpha', param_range=ccp_alphas,
    train_scores=train_error_rates, test_scores=test_error_rates, score_name="Accuracy"
)
display.plot()
plt.scale('log')
plt.ylabel ('Error Rate')
plt.show()
```

```
In [ ]: # Final the optimal alpha value with lowest error rates.
mean_test_error = np.mean(test_error_rates, axis=1)
std_test_error = np.std(test_error_rates, axis=1) # Corrected variable name
df_ccp = pd.DataFrame({
    'ccp_alpha': ccp_alphas,
    'mean_test_score': mean_test_error,
    'std_test_score': std_test_error
})
pd.set_option('display.max_rows', None)
df_ccp.sort_values('mean_test_score').head()
```

```
Out[ ]:
```

	ccp_alpha	mean_test_score	std_test_score
299	0.002056	0.387181	0.016560
298	0.002053	0.387181	0.016560
300	0.002166	0.391165	0.016627
294	0.001826	0.392609	0.015251
295	0.001839	0.392972	0.015494

```
In [ ]: # Prune the tree based on the optimal ccp (best pruned tree)
Pruned_clf = DecisionTreeClassifier(ccp_alpha=0.002056) # Replace 0.002056 with your optimal ccp_alpha value
Pruned_clf.fit(X_train, Y_train)
```

Out []:

DecisionTreeClassifier ⓘ ?

DecisionTreeClassifier(ccp_alpha=0.002056)

In []:

```
# accuracy
# model evaluation predicted class (start here)
Y_pred = Pruned_clf.predict(X_test)
dt_accuracy = accuracy_score(Y_test, Y_pred)
dt_accuracy
```


Out []:

0.6078147612156295

In []:

```
# Confusion matrix
cm = confusion_matrix(Y_test, Y_pred)
print(cm)
sns.heatmap(cm, annot=True, cmap="YlGnBu") # Corrected colormap name
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
[[178 185]
 [ 86 242]]
```

Bagging classifier

```
In [ ]: tree_clf = DecisionTreeClassifier()
bag_clf = BaggingClassifier(
    estimator = tree_clf,
    n_estimators = 100, # The number of base estimators in the ensemble
    bootstrap = True, # Whether samples are drawn with replacement. If False, sampling without replacement is performed.
    n_jobs = -1, # The number of jobs to run in parallel; -1 means using all processors
    random_state = 30)

bag_clf.fit(X_train, Y_train)
```


```
Out[ ]: > BaggingClassifier ⓘ ?
> estimator: DecisionTreeClassifier
    > DecisionTreeClassifier ⓘ
```

```
In [ ]: Y_pred = bag_clf.predict(X_test)
        bagging_accuracy = accuracy_score(Y_test, Y_pred)
        bagging_accuracy
```

```
Out[ ]: 0.6367583212735166
```

```
In [ ]: cm = confusion_matrix(Y_test, Y_pred)
        print(cm)
        sns.heatmap(cm, annot = True, cmap = "YlGnBu")
        plt.xlabel("Predicated")
        plt.ylabel("Actual")
        plt.title("Confusion Matrix")
        plt.show()
```

```
[[228 135]
 [116 212]]
```



Boostraping Classifier


```
In [ ]: ada_clf = AdaBoostClassifier (
        estimator = tree_clf,
```

```
n_estimators = 100, # The maximum number of estimators at which boosting is terminated.
algorithm = "SAMME.R",
learning_rate = 0.5, # Weight applied to each classifier at each boosting iteration.
random_state = 30)
```

```
ada_clf.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
```

Out []:



In []:

```
Y_pred = ada_clf.predict(X_test)
bootstrapping_accuracy = accuracy_score(Y_test, Y_pred)
bootstrapping_accuracy
```


Out []:

```
0.5615050651230101
```

In []:

```
cm = confusion_matrix(Y_test, Y_pred)
print(cm)
sns.heatmap(cm, annot = True, cmap = "YlGnBu")
plt.xlabel("Predicated")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
[[209 154]
 [149 179]]
```



RandomForest Classifier

```
In [ ]: X.shape
```

```
Out[ ]: (3452, 40)
```


```
In [ ]: rf_clf = RandomForestClassifier(n_estimators=50, random_state=42, max_features=2, oob_score=True) # corrected argument name
```

```
In [ ]: from sklearn.model_selection import GridSearchCV
```

```
param_grid = {'max_features': list(range(1, 41))}
rf_grid_search = GridSearchCV(estimator=rf_clf, param_grid=param_grid, cv=5, n_jobs=-1, verbose=2)
rf_grid_search.fit(X_train, Y_train)
```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

Out []:



In []:

```
best_grid_rf = rf_grid_search.best_estimator_
print(best_grid_rf)
```

```
RandomForestClassifier(max_features=36, n_estimators=50, oob_score=True,
                      random_state=42)
```

In []:

```
# Instantiate Random Forest Classifier
rf_clf2 = RandomForestClassifier(n_estimators=50, random_state=42, max_features=36)
# Train classifier
rf_clf2.fit(X_train, Y_train)
# Make predictions and evaluate
Y_pred_rf2 = rf_clf2.predict(X_test)
RandomForestClassifier_accuracy = accuracy_score(Y_test, Y_pred_rf2)
RandomForestClassifier_accuracy
```


Out []:

```
0.6121562952243126
```

In []:

```
# 2. confusion matrix
cm = confusion_matrix (Y_test, Y_pred)
print(cm)
sns.heatmap(cm, annot = True, cmap = "YlGnBu")
plt.xlabel("Predicated" )
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
[[209 154]
 [149 179]]
```



```
In [ ]: # Final Accuracy for the four classification models
# Final Accuracy for the four classification models
print('Decision Tree Classifier: ' + str(dt_accuracy))
print('Bagging Classifier: ' + str (bagging_accuracy))
print('Boostraping Classifier: ' +str (boostraping_accuracy))
print('RandomForest Classifier: ' +str (RandomForestClassifier_accuracy))
```

```
Decision Tree Classifier: 0.6078147612156295
Bagging Classifier: 0.6367583212735166
Boostraping Classifier: 0.5615050651230101
RandomForest Classifier: 0.6121562952243126
```